

InclusiveSpaces (CAT) — Database Operations Guide

Scope: Supabase Postgres + PostGIS + pgRouting operations for routing + comfort layers

0) Big picture (what lives where)

0.1 What the backend needs (hard requirement)

Per city, the API expects a **pgRouting-style graph**:

- **Edges table:** <prefix>_ways
 - required columns (minimum): gid, source, target, cost, the_geom
 - plus **many * _weight columns** used in /api/accessibility cost expression
- **Vertices table:** <prefix>_ways_vertices_pgr
 - required columns (minimum): id, the_geom (Point)

0.2 What is NOT required in DB (for frontend visualization)

Most visualization layers are loaded in the frontend from public/data/<city>/... or WMS.

1) Prerequisites (install once)

1.1 Tools (install once per machine)

Environment

- Windows 11 + WSL2 (Ubuntu 22.04+) (recommended)
- Or Linux (Ubuntu)

Required for CAT database operations (OSM + PostGIS uploads + CLI access)

- **psql (PostgreSQL client only)**
 - Used to connect to Supabase Postgres and run SQL scripts.
- **GDAL (ogr2ogr)**
 - Used to upload GeoPackage/GeoJSON layers into PostGIS.
- **osmium-tool (osmium)**
 - Used to extract city OSM data from a large .osm.pbf file and convert it to .osm.
- **osm2pgRouting**
 - Used to import .osm into pgRouting tables (<>_ways, <>_ways_vertices_pgr).

Optional (recommended for visual checks / fixing CRS & geometry)

- **QGIS**: Use it to check CRS (= EPSG:4326), visualize and inspect layers before upload.
-

1.2 Install tools

Windows + WSL2

1. Install WSL2 + Ubuntu

- Windows PowerShell (Admin):

```
wsl --install -d Ubuntu
```

- Restart if prompted, then open **Ubuntu** terminal.

2. Install the required CLI tools inside Ubuntu (WSL)

```
sudo apt update  
sudo apt install -y postgresql-client gdal-bin osmium-tool osm2pgsql
```

3. Install QGIS (recommended)

- **Preferred (simplest)**: install **QGIS on Windows** using the official installer (GUI runs smoother than inside WSL).
-

1.3 Verify installations

- Run:

```
psql --version  
ogr2ogr --version  
osmium --version  
osm2pgsql --version
```

2) Supabase connection

2.1 Get the connection parameters

```
Host: aws-0-eu-central-1.pooler.supabase.com  
Port: 6543  
Database: postgres  
User: postgres.tcxrvmwzddsyivnfurdx  
Pool mode: transaction
```

You can also check from: [Supabase Dashboard](#) →

[Connect](#)

The screenshot shows the 'Connection String' tab selected in the Supabase dashboard. A red box highlights the 'Method' dropdown which is set to 'Transaction pooler'. Below the dropdown, there's a note: 'Learn how to connect to your Postgres databases. [Read docs](#)'.

Transaction pooler (SHARED POOLER)

Ideal for stateless applications like serverless functions where each interaction with Postgres is brief and isolated.

```
postgresql://postgres.tcxrvmwzddsyivnfurdx:[YOUR-PASSWORD]@aws-0-eu-central-1.pooler.supabase.com:6543/postgres
```

Does not support PREPARE statements

View parameters

```
host: aws-0-eu-central-1.pooler.supabase.com
port: 6543
database: postgres
user: postgres.tcxrvmwzddsyivnfurdx
pool_mode: transaction
```

For security reasons, your database password is never shown.

IPv4 compatible
Transaction pooler connections are IPv4 proxied for free.

2.2 Use the Pooler connection string

Use this format (replace [YOUR-PASSWORD])

```
postgresql://postgres.tcxrvmwzddsyivnfurdx: [YOUR-PASSWORD] @aws-0-eu-central-1.pooler.supabase.com:6543/postgres
```

2.3 Quick connection test (psql)

```
psql " postgresql://postgres.tcxrvmwzddsyivnfurdx: [YOUR-PASSWORD] @aws-0-eu-central-1.pooler.supabase.com:6543/postgres "
```

[YOUR-PASSWORD]: incspace123456

4) Database initialization (run once per DB)

Run in Supabase SQL editor or psql:

```
CREATE EXTENSION IF NOT EXISTS postgis;
CREATE EXTENSION IF NOT EXISTS pgrouting;
```

Verify:

```
SELECT PostGIS_Version();
SELECT pgr_version();
```

5) Routing graph build (per city) — OSM → osmium → osm2pgRouting

This is the core of [/api/accessibility.js](#)

5.1 Inputs you must have (choose based on data availability)

Data needed

- An OSM extract in [.osm.pbf](#) or [.osm](#) that contains the city road network.

Optional: only needed if you must cut a smaller area (common for small cities / districts)

- **City boundary polygon** in GeoJSON, CRS EPSG:4326, valid Polygon/MultiPolygon.
(Used for osmium extract --polygon ...)
-

5.2 Get OSM data (pick one route)

Route A — City-level extract available (recommended if available)

- Download a [city-level .osm.pbf](#) from an extract provider (e.g., Geofabrik). ([Geofabrik Download Server](#))
Result: <city>.osm.pbf

Route B — No city-level extract (use regional + clip)

- Download a [regional .osm.pbf](#) (country/region-level) from an extract provider such as Geofabrik. ([Geofabrik Download Server](#))
Example: greece-latest.osm.pbf (for Penteli-style cases)
 - Ensure you have <city>_boundary.geojson (EPSG:4326, valid polygon).
-

5.3 Extract city OSM using osmium (ONLY for Route B)

Syntax:

```
cd /path/to/resource/data/<City>
osmium extract --polygon <city>_boundary.geojson <region>.osm.pbf -o
<city>.osm.pbf
```

5.4 Convert to .osm (do this for both Route A and Route B) :

Syntax:

```
osmium cat <city>.osm.pbf -o <city>.osm --overwrite
```

Example:

```
osmium extract --polygon penteli_boundary.geojson greece-latest.osm.pbf -o
penteli.osm.pbf
osmium cat penteli.osm.pbf -o penteli.osm --overwrite
```

Why convert to .osm? osm2pgrouting expects OSM XML; passing .osm.pbf often fails or creates empty tables.

5.5 Import using osm2pgrouting (table prefix matters)

osm2pgrouting supports --schema and --prefix ([GitHub](#)).

Use --prefix to point to cities, so tables become <prefix>_ways and <prefix>_ways_vertices_pgr:

```
osm2pgrouting \
-f <city>.osm \
-d postgres \
-U postgres.tcxrvmwzddsyivnfurdx \
-W <DB_PASSWORD> \
-h aws-0-eu-central-1.pooler.supabase.com \
-p 6543 \
--schema public \
--prefix <prefix> \
--clean
```

<DB_PASSWORD>: incspace123456

5.6 Verify routing tables immediately

In SQL:

```
SELECT COUNT(*) AS n_edges FROM <prefix>_ways;
SELECT COUNT(*) AS n_vertices FROM <prefix>_ways_vertices_pgr;

-- Check geometry SRID (expect 4326 in most CAT workflows)
SELECT ST_SRID(the_geom) FROM <prefix>_ways LIMIT 5;
SELECT ST_SRID(the_geom) FROM <prefix>_ways_vertices_pgr LIMIT 5;
```

If SRID not 4326: decide one standard and keep it consistent with your API nearest-vertex query.

5.7 Indexing (strongly recommended for speeding up) :

```
CREATE INDEX IF NOT EXISTS <prefix>_ways_geom_gist ON <prefix>_ways USING GIST  
(the_geom);  
CREATE INDEX IF NOT EXISTS <prefix>_ways_source_idx ON <prefix>_ways (source);  
CREATE INDEX IF NOT EXISTS <prefix>_ways_target_idx ON <prefix>_ways (target);  
CREATE INDEX IF NOT EXISTS <prefix>_vertices_geom_gist ON  
<prefix>_ways_vertices_pgr USING GIST (the_geom);  
  
ANALYZE <prefix>_ways;  
ANALYZE <prefix>_ways_vertices_pgr;
```

6) Upload comfort/environment layers to PostGIS

Syntax:

```
ogr2ogr -f "PostgreSQL" \  
    PG:"host=aws-0-eu-central-1.pooler.supabase.com  
    user=postgres.tcxrvmwzddsyivnfurdx dbname=postgres password=<DB_PASSWORD>  
    port=6543" \  
    /path/input.geojson \  
    -nln <target_table> \  
    -progress
```

<DB_PASSWORD>: incspace123456

7) * _weight columns on <prefix>_ways (required for comfort routing)

— schema + population

7.1 What these columns mean (IMPORTANT)

In [pages/api/accessibility.js](#), routing cost is computed like:

- If a DB column indicates the factor is “present / active”, then the SQL uses the **front-end user value** as a multiplier.
- Otherwise it uses 1 (neutral, no penalty) when calculates the reachable area.
- Whether the factor is “present / active” depends on how strong the environment factor is.
(e.g. noise level > 65db → noise_weight = 1)

So DB columns must be binary flags (0/1), it only says whether a factor applies on that edge.

There are **two types** of flags, matching the current accessibility.js CASE logic:

Type A — Negative factor (bad = 1, neutral = 0)

noise_weight, trafficLight_weight, temp_weight_s, temp_weight_w, path_width_weight,
stair_weight, obstacle_weight, slope_weight, uneven_surfaces_weight, poor_pavement_weight,
kerbs_h_weight, pedestrian_flow_weight

Type B — Positive factor (good = 1, missing/negative = 0)

light_weight, tactile_weight, tree_weight, green_weight, station_weight, wc_d_weight,
facilities_weight

7.2 Add missing columns (safe + repeatable)

Use defaults 1.0 (positive environmental factor) or 0.0 (negative environmental factor) so routing still works even if you haven't computed a factor.

```
ALTER TABLE <prefix>_ways
ADD COLUMN IF NOT EXISTS noise_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS light_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS trafficLight_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS tactile_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS tree_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS temp_weight_s double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS temp_weight_w double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS blue_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS green_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS station_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS wc_d_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS path_width_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS stair_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS obstacle_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS slope_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS uneven_surfaces_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS poor_pavement_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS kerbs_h_weight double precision NOT NULL DEFAULT 0.0,
ADD COLUMN IF NOT EXISTS facilities_weight double precision NOT NULL DEFAULT 1.0,
ADD COLUMN IF NOT EXISTS pedestrian_flow_weight double precision NOT NULL DEFAULT 0.0
```

8) Populate flags from layers (binary 0/1) — the “how”

8.1 General workflow (per factor)

1. Import the layer to DB (Step 6) → a table like streetlight, noise_distribution, obstacle, etc.
 2. Ensure layer has a geometry column geom, EPSG is 4326
 3. Add spatial index to speed up updates:

```
CREATE INDEX IF NOT EXISTS <layer>_geom_gix ON <layer> USING GIST (geom);
```
 4. Run an UPDATE <prefix>_ways ... to set the corresponding *_weight column to 0/1.
-

8.2 Pattern A — Point layers (distance-based → set 0/1)

Rule of thumb (buffer / distance):

For point layers (e.g., trees, street lights), when relating points to routing edges, a **20 m** distance threshold usually works well (i.e., treat an edge as “affected/covered” if a point is within ~20 m).

Example: **lighting** as environmental factor

(as positive factor → default=1; if no streetlight nearby, discomfort → value = 0)

```
UPDATE <prefix>_ways w
SET light_weight =
CASE
    WHEN EXISTS (
        SELECT 1
        FROM streetlight s
        WHERE ST_DWithin(w.the_geom::geography, s.geom::geography, 20)
    )
    THEN 1.0
    ELSE 0.0
END;
```

8.3 Pattern B — Polygon zones (intersect-based → set 0/1)

Rule of thumb (two polygon types):

1. Scattered patches that rarely touch the road geometry

Example: **green infrastructure** areas.

→ Use a **20 m buffer** of polygons and then test intersection.

2. Full-coverage surfaces over the whole city

Example: **noise distribution, temperature heatmaps**.

→ Prefer **threshold filtering** (e.g., dB_level >= 65, -2 < thermal_comfort_level < 2) and then set flags based on the thresholded zones.

Example: **noise** zones with attribute dB_level

(negative factor → default=0; if in a high-noise zone → discomfort, value set to 1)

```
UPDATE <prefix>_ways w
SET noise_weight =
CASE
    WHEN EXISTS (
        SELECT 0
        FROM noise_distribution n
        WHERE ST_Intersects(w.the_geom, n.geom)
            AND n. dB_level >= 65      -- threshold policy (adjust per dataset)
    )
    THEN 1.0
    ELSE 0.0
END;
```

8.4 Pattern C — Line layers (buffer then intersect → set 0/1)

Rule of thumb (buffer):

For line layers, it's recommended to apply a **20 m buffer** before checking intersection with routing edges (lines often represent a feature corridor rather than an exact road surface match)

Example: **tactile** lines buffered by 10m

(positive factor → default=1 (means tactile support exists); if no tactile support → set to 0)

```
UPDATE <prefix>_ways w
SET tactile_weight =
CASE
    WHEN EXISTS (
        SELECT 1
        FROM tactile_lines t
        WHERE ST_Intersects(
            w.the_geom,
            ST_Buffer(t.geom::geography, 20)::geometry
        )
    )
    THEN 1.0
    ELSE 0.0
END;
```

9 Cleanup (recommended: keep Supabase DB lean)

After you have successfully populated all *_weight flags on <prefix>_ways, **remove the imported environment layer tables** (keep a local backup if needed) to keep the Supabase database tidy and within the **free storage quota**.

Recommended to keep on Supabase (per city)

- <prefix>_ways
- <prefix>_ways_vertices_pgr

Then clean up storage

- Drop the environment layer tables you no longer need (e.g., noise_distribution, streetlight...).
- Run **VACUUM** (and optionally ANALYZE) to reclaim space and refresh stats:)

```
VACUUM (ANALYZE) <prefix>_ways;  
VACUUM (ANALYZE) <prefix>_ways_vertices_pgr;
```

(If VACUUM is restricted by permissions in your Supabase setup, skip it—dropping tables still saves most space.)
